

Introduction to Metaheuristics

Estimation of Distribution

DrC. Alejandro Piad Morffis

CC-BY - matcom.in/metaheuristics

Review: The metaheuristic framework

- ▶ given a black box function F
- ▶ define stop criteria
- ▶ while not stop do:
 - ▶ generate a new solution x_i (or many)
 - ▶ evaluate $y_i = F(x_i)$
 - ▶ update global best $y^* \leftarrow \min\{y^*, y_i\}$
 - ▶ learn something about F for next iteration (maybe)
- ▶ return global best y^*

Review: What can we learn...

► ... in local search?:

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima
- ▶ ... **in evolutionary search**?

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima
- ▶ ... **in evolutionary search**?
 - ▶ Components of good solutions

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima
- ▶ ... **in evolutionary search**?
 - ▶ Components of good solutions
- ▶ ... **in swarm optimization**?

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima
- ▶ ... **in evolutionary search**?
 - ▶ Components of good solutions
- ▶ ... **in swarm optimization**?
 - ▶ Features of the global landscape

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Review: What can we learn...

- ▶ ... **in local search**?:
 - ▶ Gradients
 - ▶ Features of local optima
- ▶ ... **in evolutionary search**?
 - ▶ Components of good solutions
- ▶ ... **in swarm optimization**?
 - ▶ Features of the global landscape
 - ▶ Features of local optima

Every metaheuristic learns some implicit **model** of the function to guide the search.

Can we make that model explicit?

Estimation of Distribution Algorithms

- ▶ given a black box function F
- ▶ **initialize** probabilistic model \mathfrak{F}
- ▶ define stop criteria
- ▶ while not stop do:
 - ▶ **sample** N new solutions $\{x_i\}$ from \mathfrak{F}
 - ▶ evaluate $\{y\}_i = \{F(x_i)\}$
 - ▶ update global best $y^* \leftarrow \min\{y^*, y_i \dots\}$
 - ▶ **update** model \mathfrak{F} for next iteration
- ▶ return global best y^*

Different EDAs differ in the selection of the model \mathfrak{F}

Types of models

What does the model estimates:

(The distribution of):

- ▶ the global optimum
- ▶ the local optima
- ▶ the best features
- ▶ fitness values
- ▶ ...

Types of models

What does the model estimates:

(The distribution of):

- ▶ the global optimum
- ▶ the local optima
- ▶ the best features
- ▶ fitness values
- ▶ ...

What assumptions the model encodes:

- ▶ Are components correlated?
- ▶ ...

PBIL: Genetic algorithms turn EDAs

Assuming the simplest possible model

PBIL: Genetic algorithms turn EDAs

Assuming the simplest possible model

- ▶ A fixed set of binary components i
- ▶ A binomial distribution for each component p_i
- ▶ Components are independent

PBIL: Genetic algorithms turn EDAs

Assuming the simplest possible model

- ▶ A fixed set of binary components i
- ▶ A binomial distribution for each component p_i
- ▶ Components are independent

Sample:

- ▶ For each component i , set $x_i = U(0, 1) < p_i$

Update rule:

- ▶ For each component i :
 - ▶ compare solutions where $x_i = 0$ vs $x_i = 1$
 - ▶ if most solutions where $x_i = 0$ are better, decrease p_i
 - ▶ else, increase p_i

PBIL: Genetic algorithms turn EDAs

Some implementation details:

- ▶ Use a learning rate $r \ll 1$
- ▶ Adjust learning rate to balance exploration/exploitation

PBIL: Genetic algorithms turn EDAs

Some implementation details:

- ▶ Use a learning rate $r \ll 1$
- ▶ Adjust learning rate to balance exploration/exploitation

Extension to categorical distributions

- ▶ For each component i :
 - ▶ For each value b_{ij} ,
 - ▶ Compute all solutions where $x_i = b_{ij}$
 - ▶ Increase p_{ij} proportional to their fitness
- ▶ Renormalize probabilities

PBIL: Genetic algorithms turn EDAs

Some implementation details:

- ▶ Use a learning rate $r \ll 1$
- ▶ Adjust learning rate to balance exploration/exploitation

Extension to categorical distributions

- ▶ For each component i :
 - ▶ For each value b_{ij} ,
 - ▶ Compute all solutions where $x_i = b_{ij}$
 - ▶ Increase p_{ij} proportional to their fitness
- ▶ Renormalize probabilities

Alternative formulation:

- ▶ Compute marginal distribution from K best solutions
- ▶ Interpolate model using learning rate

UMDA: Moving to continuous spaces

Model:

- ▶ Each component is a normal distribution $N(\mu_i, \sigma_i)$
- ▶ Start with large σ values

Update rule:

- ▶ Computer marginal distribution of best solutions
- ▶ Compute μ_b, σ_b
- ▶ Interpolate

Removing the independence constraint

Instead of N independent Gaussian univariate distributions, use a multivariate distribution

- ▶ A vector $\mathbf{\mu}$ of μ_i means
- ▶ A matrix Σ of σ_{ij} covariance pairs

Update rule:

- ▶ Keep the best $M < N$ solutions
- ▶ Update mean
- ▶ Update covariance matrix on solution pairs

Better idea: **CMA-ES**

CMA-ES

Updating the covariance matrix w.r.t. solutions naively is suboptimal. (Why?)

CMA-ES

Updating the covariance matrix w.r.t. solutions naively is suboptimal. (Why?)

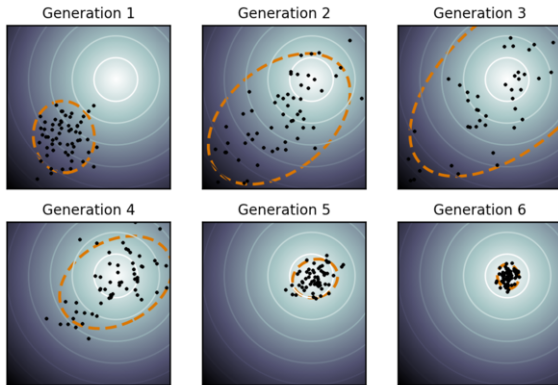
CMA-ES updates the covariance matrix such that it produces new solutions in the **direction of the best search steps** seen so far.

Details are involved:

- ▶ Compute incremental evolution path: *If the mean moves in a similar direction, the path length increases.*
- ▶ Update step-size σ in a similar way.
- ▶ Compute covariance matrix.

CMA-ES

Adapts much quicker to the function landscape.



State-of-the-art in continuous optimization.

Towards general-purpose domains

What happens when the domain is too complex?:

- ▶ Continuous and discrete parameters
- ▶ Conditional parameters (hierarchical)

Example: Optimizing in the space of computer programs.

- ▶ Analytic regression
- ▶ Compiler optimization
- ▶ AutoML

Towards general-purpose domains

What happens when the domain is too complex?:

- ▶ Continuous and discrete parameters
- ▶ Conditional parameters (hierarchical)

Example: Optimizing in the space of computer programs.

- ▶ Analytic regression
- ▶ Compiler optimization
- ▶ AutoML

Define the search space with a probabilistic grammar

Probabilistic Grammatical Evolution

Model:

- ▶ Associate probabilities to grammar productions
- ▶ Sample in DFS order

Update rule:

- ▶ Select $K < N$ best solutions
- ▶ Compute marginal probability of used productions
- ▶ Update model by interpolation

What is the fundamental problem with PGE?

You have to evaluate many solutions for a single model update. For e.g., AutoML, this can be prohibitive.

Bayesian Optimization

Can we estimate with a single solution?

Bayesian Optimization

Can we estimate with a single solution?

- ▶ Assume F can be approximated by, e.g., a mixture of Gaussians
- ▶ Bake uncertainty into the sampling process

Bayesian Optimization

Can we estimate with a single solution?

- ▶ Assume F can be approximated by, e.g., a mixture of Gaussians
- ▶ Bake uncertainty into the sampling process

Update rule:

- ▶ Build \mathfrak{F} from previous evaluations
- ▶ Construct **acquisition function** G
- ▶ Find global optimum in G

Examples of acquisition functions:

- ▶ Probability of improvement
- ▶ Expected improvement
- ▶ Confidence bounds
- ▶ ...

Some conclusions

Why estimation of distribution works:

- ▶ EDAs are some of the most robust optimization strategies
- ▶ Easy to incorporate domain knowledge (e.g., what general shape you expect the function to have)
- ▶ Some empirical and formal convergence guarantees

Other considerations:

- ▶ PBIL / UMDA are simple and good enough for most cases.
- ▶ Bayesian Optimization is the fastest when F is very costly.
- ▶ **Do not** reimplement CMA-ES! Use a library instead.
- ▶ PGE is Turing-complete :)

Time to practice!

